

Friday AI - Complete Backend Architecture: SIP → LiveKit → Agent

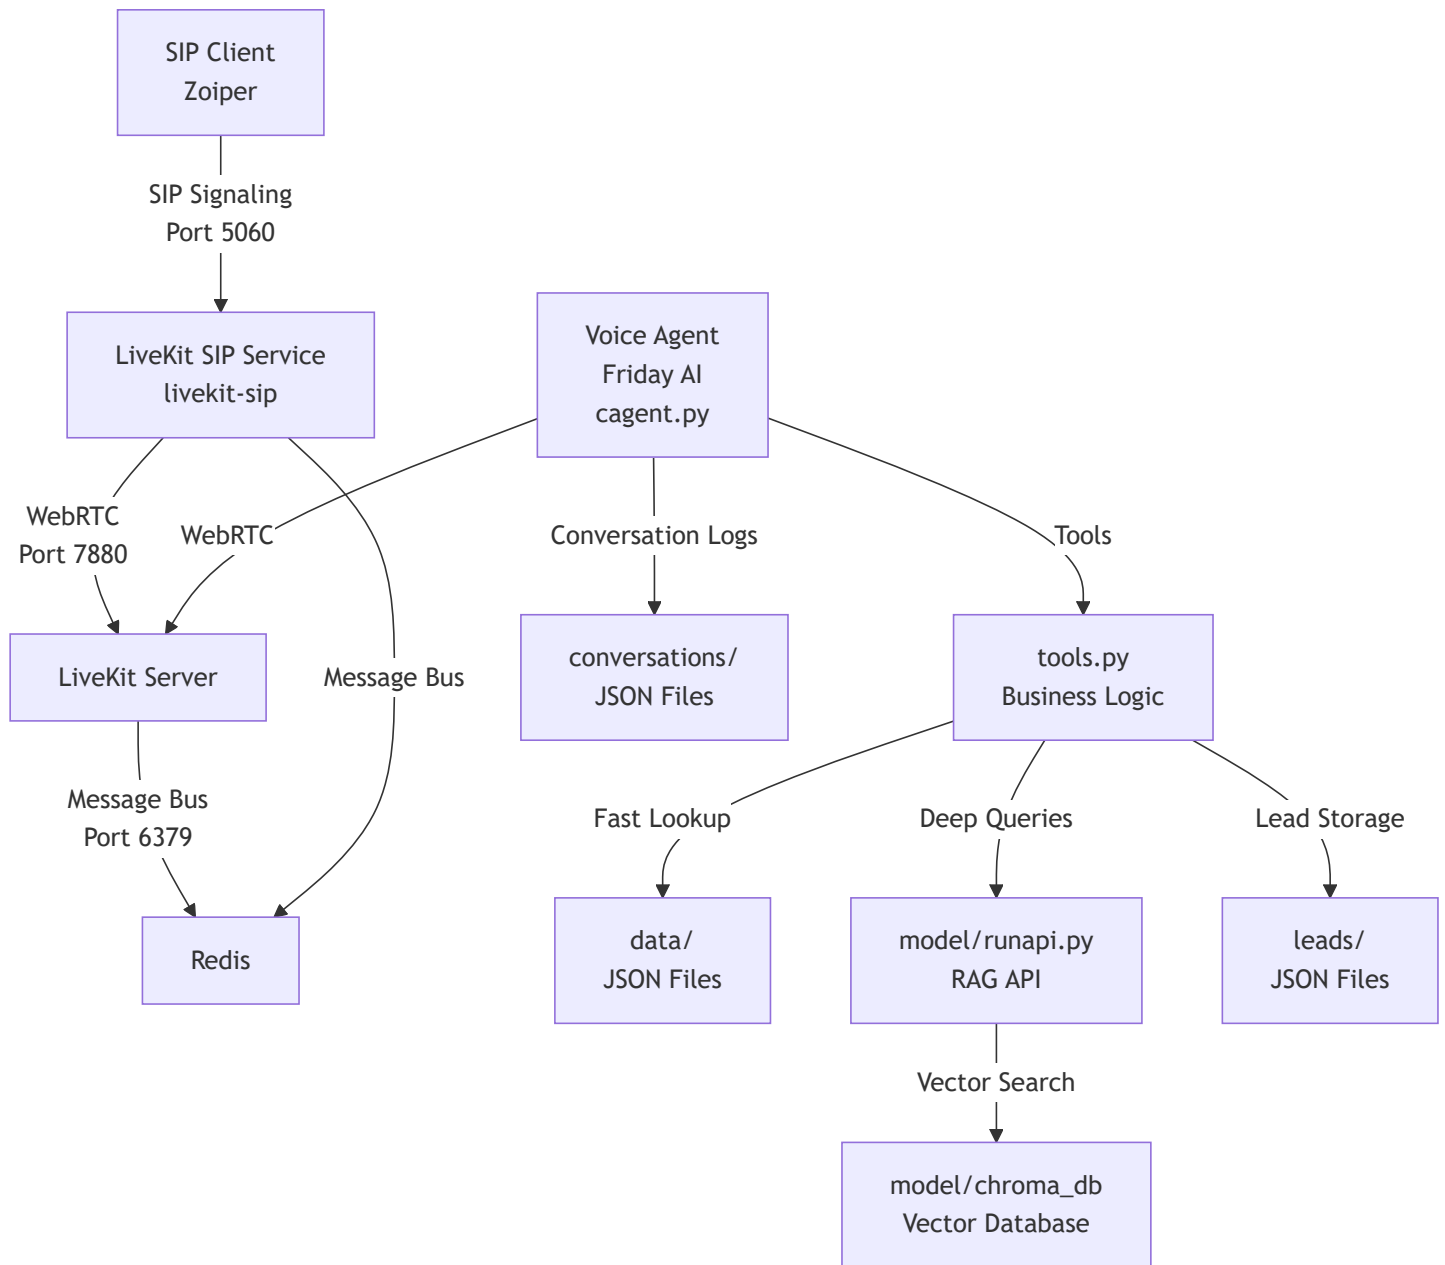
This document describes the complete Friday AI backend architecture including SIP telephony integration, LiveKit real-time communication, voice agent processing, and data flows. It serves as the definitive reference for developers working on the Friday AI system.

1. High-level Architecture Overview

Friday AI is a complete voice assistant system with SIP telephony integration, featuring:

- **SIP Telephony:** Zoiper/softphone → LiveKit SIP bridge → LiveKit server → Voice agent
- **Hybrid Knowledge:** Fast JSON lookup + RAG-powered vector database for complex queries
- **Lead Capture:** Automatic lead detection, validation, and storage in Hinglish conversations
- **Real-time Voice:** STT → Agent processing → TTS with LiveKit WebRTC infrastructure
- **Conversation Logging:** Persistent conversation history with timestamps and lead tracking

Core Components

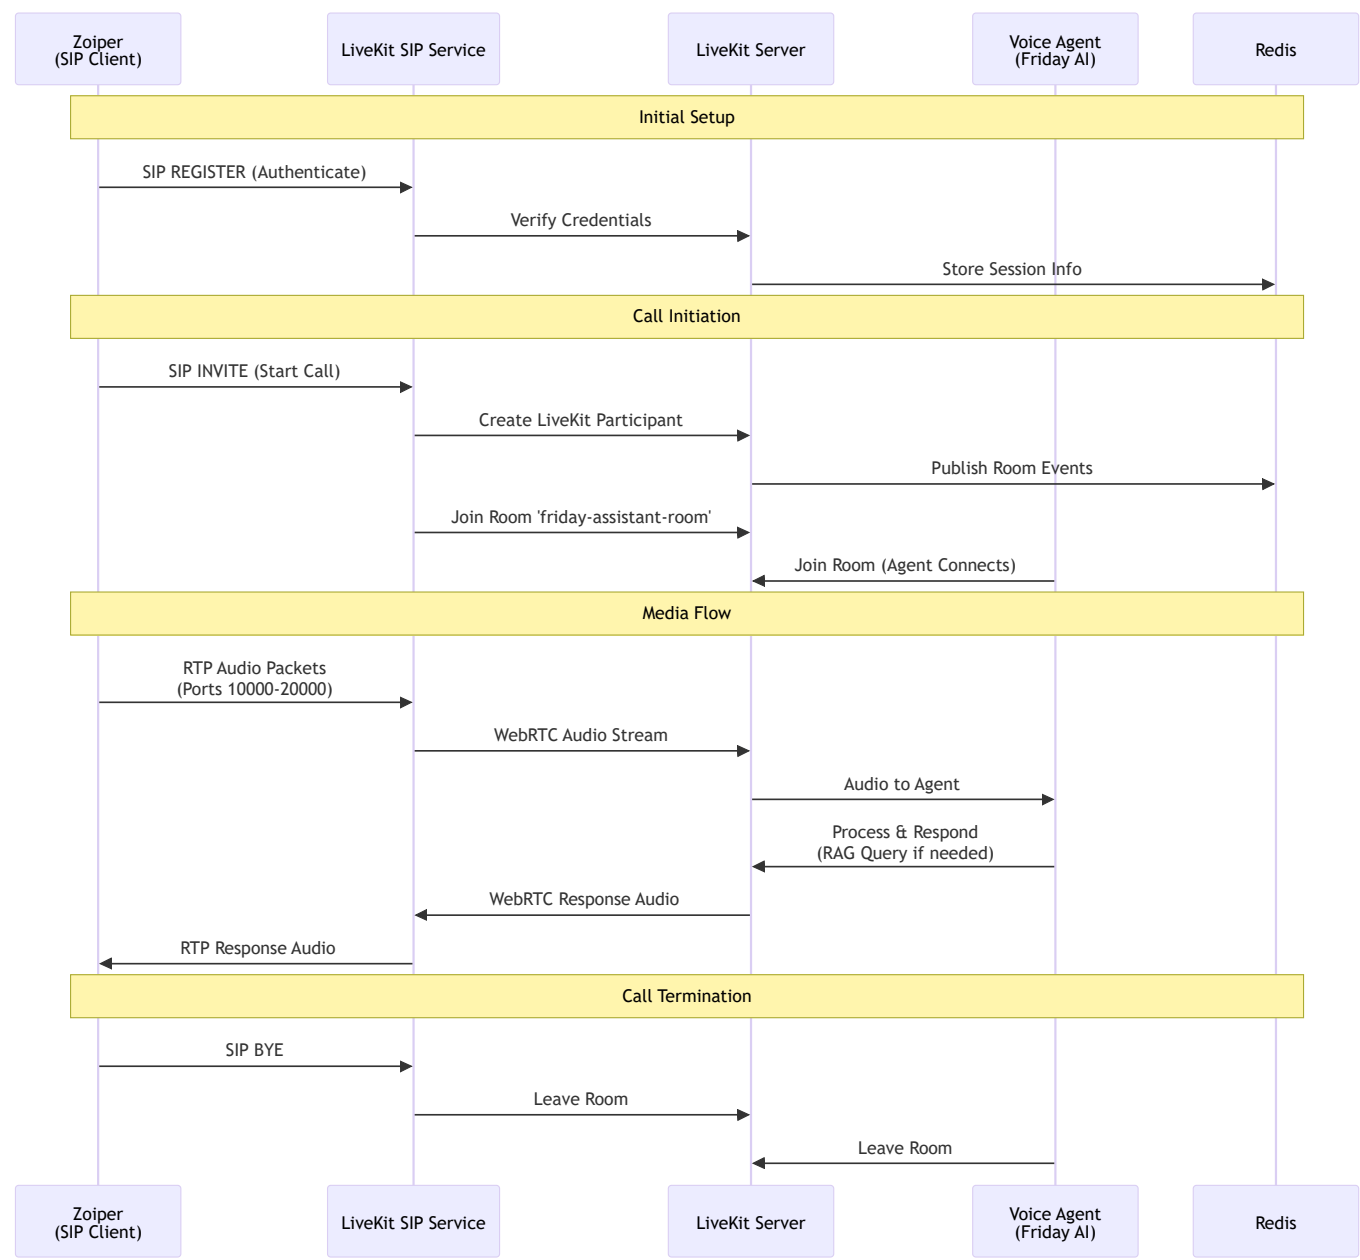


Service Architecture

- **LiveKit Server:** Core WebRTC service managing rooms, participants, and media tracks
- **LiveKit SIP Service:** SIP-to-WebRTC bridge handling phone call signaling and RTP conversion
- **Redis:** Message bus enabling communication between LiveKit server and SIP service
- **Voice Agent:** Python agent (`cagent.py`) joining LiveKit rooms for conversation processing
- **RAG System:** Vector database with embeddings for knowledge retrieval
- **Lead Management:** Automatic lead detection and JSON storage system

2. SIP Telephony Integration

SIP Call Flow



Configuration Requirements

LiveKit Server (`livekit.yaml`):

```

port: 7880
bind_addresses:
  - 0.0.0.0
keys:
  APIntavBoHTqApw: pRkd16t4uYVUs9nSlNeMawSE1qmUzfV2ZkSrMT2aiFM
redis:
  address: '192.168.109.66:6379'
  db: 0

```

SIP Service (sip-setup/config.yaml):

```

api_key: APIntavBoHTqApw
api_secret: pRkd16t4uYVUs9nSlNeMawSE1qmUzfV2ZkSrMT2aiFM
ws_url: ws://192.168.109.66:7880
redis:
  address: 192.168.109.66:6379
sip_port: 5060
rtp_port: 10000-20000
use_external_ip: false
logging:
  level: debug

```

Critical Setup Commands:

```

# Start services in order
redis-cli ping                                # verify Redis
livekit-server --config livekit.yaml          # start main server
cd sip-setup && ./livekit-sip --config config.yaml # start SIP bridge

# Setup SIP trunk and dispatch
lk project add friday --url ws://192.168.109.66:7880 --api-key APIntavBoHTqApw --api-secret pRkd16t4
lk sip inbound create --project friday inbound_trunk.json
lk sip dispatch create --project friday sip_dispatch.json

# Start voice agent
python cagent.py

```

Automating SIP Trunk & Dispatch creation

To avoid manual copy/paste when creating SIP trunks and dispatch rules, use the `lk` CLI and parse its JSON output with `jq`. Example (run from repository root where `sip-setup/` lives):

```
# Create inbound trunk and capture its ID
TRUNK_ID=$(lk sip inbound create --project friday sip-setup/inbound_trunk.json | jq -r '.sip_trunk_id')

# Replace placeholder in dispatch JSON and create the dispatch rule
sed -i "s/REPLACE_WITH_TRUNK_ID/$TRUNK_ID/g" sip-setup/sip_dispatch.json
lk sip dispatch create --project friday sip-setup/sip_dispatch.json
```

This follows the canonical pattern in `README.md` and prevents human error during provisioning.

Start services (development)

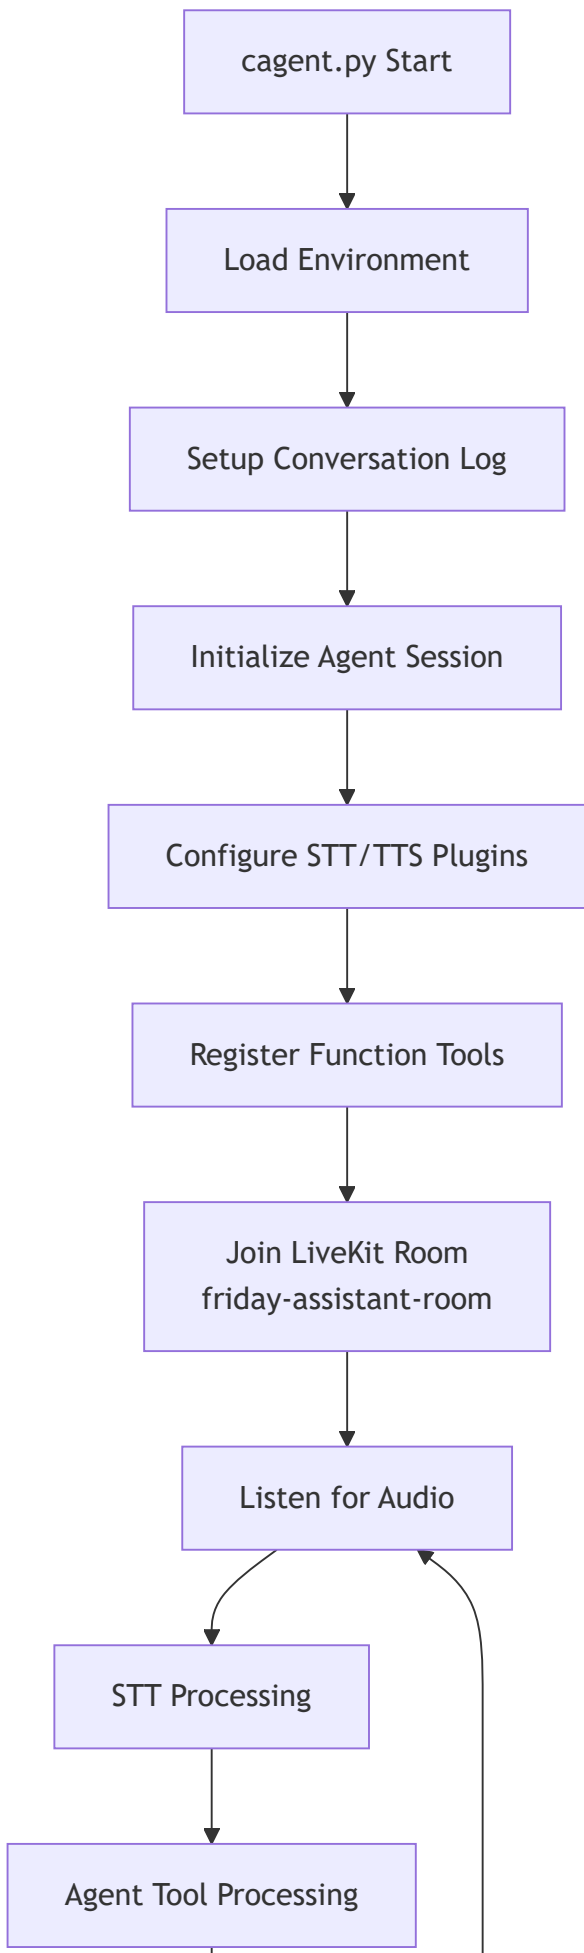
For local development, use `screen` to run services detached. For production prefer `systemd` or container orchestration.


```
screen -dmS livekit-server livekit-server --config sip-setup/livekit.yaml
screen -dmS sip-bridge livekit-sip --config sip-setup/config.yaml
screen -dmS friday-agent bash -c "source ainvenv/bin/activate && python cagent.py"
```

3. Voice Agent Core ([cagent.py](#))

The voice agent is the central component that connects to LiveKit rooms and processes conversations.

Agent Session Flow





Key Components

- **Agent Session:** Manages LiveKit connection, room participation, and audio processing
- **STT/TTS Plugins:** Deepgram for speech-to-text, Cartesia for text-to-speech (Hindi voice)
- **Function Tools:** Business logic tools registered with `@function_tool()` decorator
- **Conversation Logging:** Realtime events are appended to `conversations/transcripts.jsonl`, and a final session snapshot is written to `conversations/transcript_session_<timestamp>.json`

Agent Configuration

```
# Agent initialization in cagent.py
session = AgentSession(
    stt=deepgram.STT(model="nova-3", language="multi"),
    llm=google.LLM(model="gemini-2.5-flash", temperature=0.8),
    tts=cartesia.TTS(
        model="sonic-2",
        language="hi",
        voice="f91ab3e6-5071-4e15-b016-cde6f2bcd222",
    ),
    vad=silero.VAD.load(),
)
```

4. Business Logic Tools ([tools.py](#))

Core business logic functions exposed to the agent via `@function_tool()` decorator pattern.

Tool Functions

`triotech_info(query: str) → str`

- **Purpose:** Handles product and company information queries
- **Flow:** Fast JSON lookup in `data/` → RAG fallback for complex queries

- **Implementation:**

- Check keywords: "features", "how to", "api", "integrate" → use RAG
- Simple product queries → return JSON snippet
- Fallback → call `model/runapi.py` for vector search


detect_lead_intent(text: str) → dict

- **Purpose:** Detects if user message indicates sales interest
- **Output:** {is_lead: bool, confidence: float, extracted_fields: dict}
- **Triggers:** Keywords like "interested", "price", "demo", "sales team"

create_lead(name, email, company, interest, phone=None, budget=None) → str

- **Purpose:** Validates and stores lead information
- **Validation:** Required fields present, email format check
- **Storage:** `leads/lead_YYYYMMDD_HHMMSS.json` with English keys
- **Response:** Hinglish confirmation message

Tool Integration Flow



5. RAG System Architecture (model/)

Hybrid knowledge retrieval system combining fast JSON lookups with vector database search.

RAG Components


Vector Database Build (`model/build_db.py`):

- Chunks knowledge documents (chunk_size=1500, overlap=200)
- Uses HuggingFace embeddings for vector generation
- Builds persistent Chroma DB at `model/chroma_db/`
- Sources: `data/triotech_knowledge.txt` and other knowledge files

RAG API Runtime (`model/runapi.py`):

- Flask API server for RAG queries (port 5000)
- Handles vector similarity search
- LLM integration with API key rotation
- Fallback handling for `ResourceExhausted` errors

RAG Query Flow



Knowledge Sources

- data/triotech_content.json : Fast lookup for basic product info
- data/triotech_knowledge.txt : Deep knowledge for RAG processing
- Vector embeddings optimized for technical queries about features, APIs, integration

API Key Management

- Automatic rotation on ResourceExhausted errors
- Environment variable configuration
- Fallback to cached responses on API failures

6. Data Models and Storage

Conversation Logging

Format: conversations/conversation_YYYYMMDD_HHMMSS.json

```
{
  "conversation": [
    {
      "role": "user|agent",
      "content": "message text",
      "timestamp": "2025-10-07T16:42:56.532626",
      "source": "google_llm|cartesia_tts"
    }
  ]
}
```

Lead Storage

Format: leads/lead_YYYYMMDD_HHMMSS.json

```
{
  "timestamp": "2025-10-07T16:46:08.546889",
  "source": "Friday AI Assistant",
  "status": "new",
  "name": "Full Name",
  "email": "email@example.com",
  "company": "Company Name",
  "interest": "Product/Service",
  "phone": "phone number",
  "budget": "budget range",
  "job_title": "",
  "timeline": ""
}
```

Key Conventions:

- User-facing language: Hinglish (Hindi + English mix)
- Lead JSON fields: English only (for CRM integration)
- Timestamps: ISO 8601 format
- File naming: YYYYMMDD_HHMMSS pattern for chronological sorting

Configuration Files

Environment Variables (.env):

- GOOGLE_API_KEY : For Gemini LLM
- HUGGINGFACE_API_KEY : For embeddings
- LIVEKIT_API_KEY / LIVEKIT_API_SECRET : For LiveKit tokens
- LLM_MODEL : Configurable model (default: "gemini-2.5-flash")

SIP Configuration (sip-setup/inbound_trunk.json):

```
{
  "trunk": {
    "name": "Zoiper Local Inbound",
    "auth_username": "1001",
    "auth_password": "1001"
  }
}
```

SIP Dispatch (sip-setup/sip_dispatch.json):

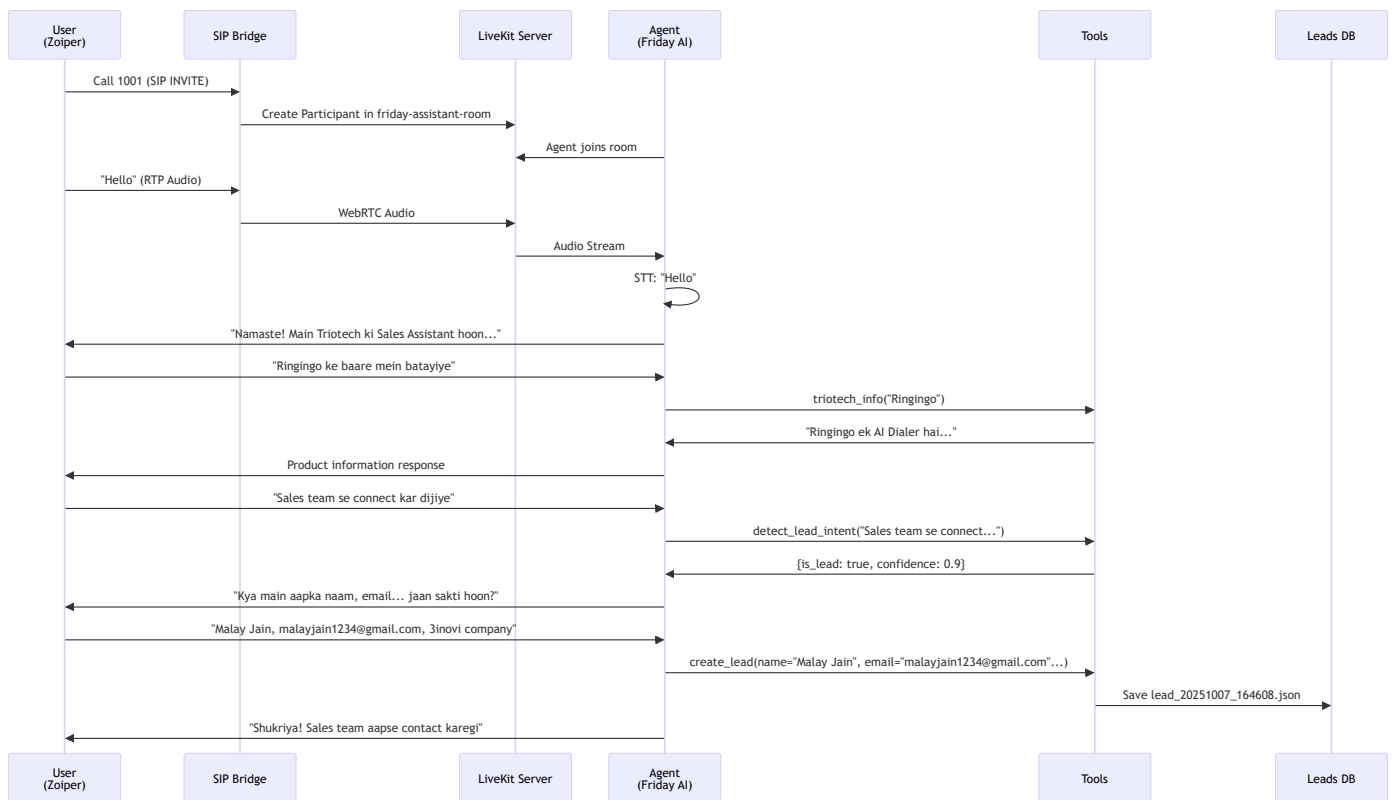
```

{
  "dispatch_rule": {
    "name": "Zoiper to Assistant Room Rule",
    "trunk_ids": ["ST_TRUNK_ID"],
    "rule": {
      "dispatchRuleDirect": {
        "roomName": "friday-assistant-room"
      }
    }
  }
}

```

7. Complete End-to-End Flow Example

Real Call Session Example



Directory Structure After Operation

Friday - Copy/

```
├─ conversations/
│   ├─ transcripts.jsonl                # Streaming events (STT chunks + committed items)
│   └─ transcript_session_2025-10-07T16-42-56.json # Complete call snapshot
├─ leads/
│   └─ lead_20251007_164608.json        # Extracted lead data
├─ sip-setup/
│   ├─ config.yaml                    # SIP bridge config
│   ├─ inbound_trunk.json             # SIP trunk definition
│   └─ sip_dispatch.json              # Call routing rules
├─ model/
│   ├─ chroma_db/                    # Vector database
│   ├─ build_db.py                   # DB builder
│   └─ runapi.py                     # RAG API server
├─ data/
│   ├─ triotech_content.json          # Fast lookup data
│   └─ triotech_knowledge.txt         # RAG knowledge base
├─ cagent.py                         # Voice agent entry point
├─ tools.py                          # Business logic functions
├─ prompts.py                        # Hinglish prompts
├─ config.py                         # Config helpers
└─ livekit.yaml                      # LiveKit server config
```


8. Deployment and Operations

Service Startup Sequence

```
# 1. Verify prerequisites
redis-cli ping                                # Should return PONG

# 2. Start LiveKit server
livekit-server --config livekit.yaml          # Port 7880

# 3. Start SIP bridge (separate terminal)
cd sip-setup && ./livekit-sip --config config.yaml # Port 5060

# 4. Setup SIP routing (one-time)
lk project add friday --url ws://192.168.109.66:7880 --api-key APIntavBoHTqApw --api-secret pRkd16t4
lk sip inbound create --project friday inbound_trunk.json
lk sip dispatch create --project friday sip_dispatch.json

# 5. Start voice agent
python cagent.py                              # Joins friday-assistant-room

# 6. Start RAG API (optional, for complex queries)
python model/runapi.py                        # Port 5000
```

Debugging and Monitoring

SIP Traffic Analysis:

```
# Monitor SIP packets
sngrep

# Check port status
ss -tulnp | grep -E "5060|7880|6379"

# Verify room participants
lk room participants --room friday-assistant-room
```

Log Analysis:

- SIP Bridge: Look for "SIP invite authentication successful"
- LiveKit Server: Check for participant join/leave events
- Agent: Monitor conversation logs in `conversations/`

- RAG API: Check vector database queries and LLM responses

Common Issues:

- **"sip not connected"**: Redis not accessible to both LiveKit server and SIP bridge
- **API key mismatch**: `livekit.yaml` and `sip-setup/config.yaml` must have identical keys
- **RTP audio issues**: Firewall blocking UDP ports 10000-20000
- **Agent not responding**: Verify agent joined correct room name

Security Considerations

- **API Secrets**: Never commit LiveKit API secrets to version control
- **PII Protection**: Leads and conversations contain sensitive data
- **Network Security**: Use firewalls to restrict SIP/RTP ports to trusted networks
- **Token Management**: Prefer backend token minting over frontend secret exposure

Performance Optimization

- **RAG Database**: Rebuild after knowledge base updates: `python model/build_db.py`
- **Conversation Cleanup**: Archive old conversations to prevent disk bloat
- **Lead Deduplication**: Implement lead validation to prevent duplicates
- **Plugin Caching**: Use modified plugins in `backup_plugin_modifications/` for consistent performance

9. Testing and Validation

Unit Tests

```
# Core functionality tests
python test_triotech_assistant.py          # Tool function tests
python test_dummy_plugins.py               # Plugin integration tests
python test_lead_detection.py              # Lead capture validation

# Build and test RAG system
python model/build_db.py                   # Rebuild vector database
python -c "import requests; print(requests.get('http://localhost:5000/health').json())"
```

Integration Tests

SIP Call Testing:

1. Configure Zoiper with SIP server `192.168.109.66:5060`

2. Register as user 1001 , password 1001
3. Dial any number to test call routing
4. Verify agent response and conversation logging

Expected Test Results:

- SIP registration successful in Zoiper
- Call connects and agent responds in Hinglish
- Conversation saved to conversations/
- Lead detection and capture working for sales inquiries
- RAG responses for complex technical queries

Quality Gates

Pre-deployment Checklist:

- ☐ All unit tests passing
- ☐ SIP bridge can authenticate test calls
- ☐ Agent joins room and responds to audio
- ☐ Lead creation saves valid JSON with English keys
- ☐ RAG API returns relevant responses
- ☐ Conversation logs are properly formatted
- ☐ No API secrets committed to version control

Development Workflow

```
# Standard development cycle
git pull origin main
pip install -r requirements.txt
python model/build_db.py           # If knowledge changed
python test_triotech_assistant.py  # Verify core functions
python cagent.py                   # Test agent locally
```

10. Developer Onboarding

Quick Start for New Developers

1. Read these files first:

- cagent.py - Agent entry point and session management
- tools.py - Core business logic and function tools
- prompts.py - Hinglish prompts and conversation rules

- Comprehensive Setup Guide_ Integrating a SIP Client with a Self-Hosted LiveKit Environment.md

2. Set up development environment:

```
# Clone and install
git clone <repository-url>
cd friday-ai
cp .env.example .env          # Fill in API keys
pip install -r requirements.txt

# Build knowledge base
python model/build_db.py

# Test basic functionality
python test_triotech_assistant.py
```

3. Understand the data flow:

- SIP calls → LiveKit rooms → Agent processing → Tools → Storage
- User queries → JSON lookup or RAG → Hinglish responses
- Lead intent → Validation → JSON storage with English keys

4. Key conventions to follow:

- User-facing strings in Hinglish
- Lead JSON keys in English only
- Function tools use `@function_tool()` decorator
- Conversation logs: `conversation_YYYYMMDD_HHMMSS.json`
- Lead files: `lead_YYYYMMDD_HHMMSS.json`

Adding New Features

1. **New tool functions:** Add to `tools.py` with decorator
2. **New prompts:** Update `prompts.py` with Hinglish templates
3. **Knowledge updates:** Edit `data/triotech_knowledge.txt` then rebuild: `python model/build_db.py`
4. **Plugin modifications:** Use `backup_plugin_modifications/` and `docker_scripts/`

Documentation Updates

When making changes, update:

- This architecture document for structural changes
- `.github/copilot-instructions.md` for quick agent guidance
- Function docstrings for API changes
- Test files for new functionality validation